



**Modified version** of the presentation for use with the textbook  
**Data Structures and Algorithms in Java, 6<sup>th</sup> edition**, by M. T.  
Goodrich, R. Tamassia, and M. H. Goldwasser, Wiley, 2014

# (2,4) Trees




# Multi-Way Search Tree

- ◆ A multi-way search tree is an ordered tree such that
  - Each internal node has at least two children and stores  $d - 1$  key-element items  $(k_i, o_i)$ , where  $d$  is the number of children
  - For a node with children  $v_1 v_2 \dots v_d$  storing keys  $k_1 k_2 \dots k_{d-1}$ 
    - ◆ keys in the subtree of  $v_1$  are less than  $k_1$
    - ◆ keys in the subtree of  $v_i$  are between  $k_{i-1}$  and  $k_i$  ( $i = 2, \dots, d - 1$ )
    - ◆ keys in the subtree of  $v_d$  are greater than  $k_{d-1}$
  - The leaves store no items and serve as placeholders




# Multi-Way Inorder Traversal

- ◆ We can extend the notion of inorder traversal from binary trees to multi-way search trees
- ◆ Namely, we visit item  $(k_i, o_i)$  of node  $v$  between the recursive traversals of the subtrees of  $v$  rooted at children  $v_i$  and  $v_{i+1}$
- ◆ An inorder traversal of a multi-way search tree visits the keys in increasing order




# Multi-Way Searching

- ◆ Similar to search in a binary search tree
- ◆ A each internal node with children  $v_1 v_2 \dots v_d$  and keys  $k_1 k_2 \dots k_{d-1}$ 
  - $k = k_i$  ( $i = 1, \dots, d - 1$ ): the search terminates successfully
  - $k < k_1$ : we continue the search in child  $v_1$
  - $k_{i-1} < k < k_i$  ( $i = 2, \dots, d - 1$ ): we continue the search in child  $v_i$
  - $k > k_{d-1}$ : we continue the search in child  $v_d$
- ◆ Reaching an external node terminates the search unsuccessfully
- ◆ Example: search for 30



# (2,4) Trees

- ◆ A (2,4) tree (also called 2-4 tree or 2-3-4 tree) is a multi-way search with the following properties
  - **Node-Size Property:** every internal node has at most four children
  - **Depth Property:** all the external nodes have the same depth
- ◆ Depending on the number of children, an internal node of a (2,4) tree is called a 2-node, 3-node or 4-node



# Height of a (2,4) Tree

◆ **Theorem:** A (2,4) tree storing  $n$  items has height  $O(\log n)$


Proof:

- Let  $h$  be the height of a (2,4) tree with  $n$  items
- Since there are at least  $2^i$  items at depth  $i = 0, \dots, h-1$  and no items at depth  $h$ , we have

$$n \geq 1 + 2 + 4 + \dots + 2^{h-1} = 2^h - 1$$


- Thus,  $h \leq \log(n + 1)$

◆ Searching in a (2,4) tree with  $n$  items takes  $O(\log n)$  time




# Insertion

- ◆ We insert a new item  $(k, o)$  at the parent  $v$  of the leaf reached by searching for  $k$ 
  - We preserve the depth property but
  - We may cause an **overflow** (i.e., node  $v$  may become a 5-node)
- ◆ Example: inserting key 30 causes an overflow



# Overflow and Split

- ◆ We handle an **overflow** at a 5-node  $v$  with a **split operation**:
  - let  $v_1 \dots v_5$  be the children of  $v$  and  $k_1 \dots k_4$  be the keys of  $v$
  - node  $v$  is replaced by nodes  $v'$  and  $v''$ 
    - ◆  $v'$  is a 3-node with keys  $k_1 k_2$  and children  $v_1 v_2 v_3$
    - ◆  $v''$  is a 2-node with key  $k_4$  and children  $v_4 v_5$
  - key  $k_3$  is inserted into the parent  $u$  of  $v$  (a new root may be created)
- ◆ The overflow may propagate to the parent node  $u$



See Figure 11.18 in the book for a parent overflow example



# Analysis of Insertion


## Algorithm *put(k, o)*

1. We search for key  $k$  to locate the insertion node  $v$
2. We add the new entry  $(k, o)$  at node  $v$
3. **while** *overflow*( $v$ )  
    **if** *isRoot*( $v$ )  
        create a new empty root  
        above  $v$   
     $v \leftarrow \text{split}(v)$

- ◆ Let  $T$  be a (2,4) tree with  $n$  items
  - Tree  $T$  has  $O(\log n)$  height
  - Step 1 takes  $O(\log n)$  time because we visit  $O(\log n)$  nodes
  - Step 2 takes  $O(1)$  time
  - Step 3 takes  $O(\log n)$  time because each split takes  $O(1)$  time and we perform  $O(\log n)$  splits
- ◆ Thus, an insertion in a (2,4) tree takes  $O(\log n)$  time


# Deletion

- ◆ We reduce deletion of an entry to the case where the item is at the node with leaf children
- ◆ We replace the entry with its inorder successor (or, equivalently, with its inorder predecessor) and delete the latter entry
- ◆ Example: to delete key 24, we replace it with 27 (inorder successor)




# Underflow and Fusion

- ◆ Deleting an entry from a node  $v$  may cause an **underflow**, where node  $v$  becomes a 1-node with one child and no keys
- ◆ To handle an underflow at node  $v$  with parent  $u$ , we consider two cases
- ◆ **Case 1:** the adjacent siblings of  $v$  are 2-nodes
  - **Fusion operation:** we merge  $v$  with an adjacent sibling  $w$  and move an entry from  $u$  to the merged node  $v'$
  - After a fusion, the underflow may propagate to the parent  $u$



# Underflow and Transfer

- ◆ To handle an underflow at node  $v$  with parent  $u$ , we consider two cases
- ◆ **Case 2:** an adjacent sibling  $w$  of  $v$  is a 3-node or a 4-node
  - **Transfer operation:**
    1. we move a child of  $w$  to  $v$
    2. we move an item from  $u$  to  $v$
    3. we move an item from  $w$  to  $u$
  - After a transfer, no underflow occurs




# Analysis of Deletion

- ◆ Let  $T$  be a  $(2,4)$  tree with  $n$  items
  - Tree  $T$  has  $O(\log n)$  height
- ◆ In a deletion operation:
  - We visit  $O(\log n)$  nodes to locate the node from which to delete the entry
  - We handle an underflow with a series of  $O(\log n)$  fusions, followed by at most one transfer
    - Each fusion and transfer takes  $O(1)$  time
- ◆ Thus, deleting an item from a  $(2,4)$  tree takes  $O(\log n)$  time


**Modified version** of the presentation for use with the textbook  
**Data Structures and Algorithms in Java, 6<sup>th</sup> edition**, by M. T.  
Goodrich, R. Tamassia, and M. H. Goldwasser, Wiley, 2014

# Red-Black Trees




# From (2,4) to Red-Black Trees

- ◆ A red-black tree is a representation of a (2,4) tree by means of a binary tree whose nodes are colored **red** or **black**
- ◆ In comparison with its associated (2,4) tree, a red-black tree has
  - same logarithmic time performance
  - simpler implementation with a single node type



# Red-Black Trees

- ◆ A red-black tree can also be defined as a binary search tree that satisfies the following properties:
  - **Root Property:** the root is black
  - **External Property:** every leaf is black
  - **Internal Property:** the children of a red node are black
  - **Depth Property:** all the leaves have the same black depth





# Height of a Red-Black Tree


◆ **Theorem:** A red-black tree storing  $n$  items has height  $O(\log n)$

Proof:

- The height of a red-black tree is at most twice the height of its associated (2,4) tree, which is  $O(\log n)$
- ◆ The search algorithm for a binary search tree is the same as that for a binary search tree
- ◆ By the above theorem, searching in a red-black tree takes  $O(\log n)$  time

# Insertion

- ◆ To insert  $(k, o)$ , we execute the insertion algorithm for binary search trees and color the newly inserted node **red**  $z$  unless it is the root
  - We preserve the root, external, and depth properties
  - If the parent  $v$  of  $z$  is black, we also preserve the internal property and we are done
  - Else ( $v$  is red ) we have a **double red** (i.e., a violation of the internal property), which requires a reorganization of the tree
- ◆ Example where the insertion of 4 causes a double red:




# Remedying a Double Red

- ◆ Consider a double red with child  $z$  and parent  $v$ , and let  $w$  be the sibling of  $v$


## Case 1: $w$ is black

- The double red is an incorrect replacement of a 4-node
- **Restructuring**: we change the 4-node replacement




## Case 2: $w$ is red

- The double red corresponds to an overflow
- **Recoloring**: we perform the equivalent of a **split**




# Restructuring

- ◆ A restructuring remedies a child-parent double red when the parent red node has a black sibling
- ◆ It is equivalent to restoring the correct replacement of a 4-node
- ◆ The internal property is restored and the other properties are preserved




# Restructuring

- ◆ There are four restructuring configurations depending on whether the double red nodes are left or right children



# Recoloring

- ◆ A recoloring remedies a child-parent double red when the parent red node has a red sibling
- ◆ The parent  $v$  and its sibling  $w$  become black and the grandparent  $u$  becomes red, unless it is the root
- ◆ It is equivalent to performing a split on a 5-node
- ◆ The double red violation may propagate to the grandparent  $u$



There is no need to re-color the root since a black node can have red and black children. Furthermore, the root is traversed in any case so the depth property is preserved

# Analysis of Insertion


## Algorithm *insert*( $k, o$ )

1. We search for key  $k$  to locate the insertion node  $z$
2. We add the new entry  $(k, o)$  at node  $z$  and color  $z$  red
3. **while** *doubleRed*( $z$ )  
    **if** *isBlack*(*sibling*(*parent*( $z$ )))  
         $z \leftarrow \text{restructure}(z)$   
    **return**  
    **else** { *sibling*(*parent*( $z$ )) is red }  
         $z \leftarrow \text{recolor}(z)$

- ◆ Recall that a red-black tree has  $O(\log n)$  height
- ◆ Step 1 takes  $O(\log n)$  time because we visit  $O(\log n)$  nodes
- ◆ Step 2 takes  $O(1)$  time
- ◆ Step 3 takes  $O(\log n)$  time because we perform
  - $O(\log n)$  recolorings, each taking  $O(1)$  time, and
  - at most one restructuring taking  $O(1)$  time
- ◆ Thus, an insertion in a red-black tree takes  $O(\log n)$  time

# Deletion

- ◆ To perform operation **remove**( $k$ ), we first execute the deletion algorithm for binary search trees
- ◆ Let  $v$  be the internal node removed,  $w$  the external node removed, and  $r$  the sibling of  $w$ 
  - If either  $v$  or  $r$  was red, we color  $r$  black and we are done
  - Else ( $v$  and  $r$  were both black) we color  $r$  **double black**, which is a violation of the depth property requiring a reorganization of the tree
- ◆ Example where the deletion of 8 causes a double black:





# Remedying a Double Black

- ◆ The algorithm for remedying a double black node  $w$  with sibling  $y$  considers three cases

**Case 1:**  $y$  is black and has a red child

- We perform a **restructuring**, equivalent to a **transfer**, and we are done (figure 11.29)

**Case 2:**  $y$  is black and its children are both black

- We perform a **recoloring**, equivalent to a **fusion**, which may propagate up the double black violation (figure 11.30)

**Case 3:**  $y$  is red

- We perform an **adjustment**, equivalent to choosing a different representation of a 3-node, after which either Case 1 or Case 2 applies (figure 11.31)

- ◆ Deletion in a red-black tree takes  $O(\log n)$  time

# Red-Black Tree Reorganization

<b>Insertion</b> remedy double red		
Red-black tree action	(2,4) tree action	result
restructuring	change of 4-node representation	double red removed
recoloring	split	double red removed or propagated up

<b>Deletion</b> remedy double black		
Red-black tree action	(2,4) tree action	result
restructuring	transfer	double black removed
recoloring	fusion	double black removed or propagated up
adjustment	change of 3-node representation	restructuring or recoloring follows

# Notes

- ◆ RB Tree Visualization:

<http://www.cs.usfca.edu/~galles/visualization/RedBlack.html>

- ◆ 2-4 Trees and RB Trees go hand in hand

- ◆ I suggest you go over chapters 11.4 and 11.5 together

- Figures are especially helpful!
- The book also has source code

- ◆ In practice, most of the time RB trees are used instead of 2-4 trees.

- ◆ But why use RB Trees if we have AVL Trees?

# AVL vs RB Trees

	Height	Search	Insert	Delete
AVL	$\sim 1.44 \log n$	$O(\log n)$	Search + Rebalance $O(\log n) + O(1)$	Search + Rebalance $O(\log n) + O(\log n)$
RB	$\sim 2 \log n$	$O(\log n)$	Search + Recolor* + Restructure $O(\log n) + O(\log n) + O(1)$	Search + Recolor* + Restructure $O(\log n) + O(\log n) + O(1)$

- The worst case recoloring is  $O(\log n)$  but it is actually **amortized  $O(1)$** ! Even though deletion is still  $O(\log n)$  in both cases, RB trees run faster. Furthermore, recoloring has a much smaller constant than tree rotations. Thus RB trees are preferred for removal intensive applications (which also implies insertion intensive since you cannot delete non-existent nodes)
- Even though both heights are  $O(\log n)$ , AVL trees are shorter making searching faster thus they should be used for lookup intensive applications OR when  $n$  gets very large to such that searching dominates the complexity